

Intelligence = Sampling and Filtering

Presented by:
John Tan Chong Min

Drawing on Lessons From:

AlphaCode – DeepMind (Li et al, 2022)

GPT-4o on ARC (Ryan Greenblatt, 2024)

Learning, Fast and Slow (John and Motani, 2023)

LLMs as a System of Multiple Expert Agents to solve ARC (John and Motani, 2023)

AlphaGo / AlphaZero (Silver et al, 2017)


LLAMA-3 8B with MCTS (Zhang et al, 2024)



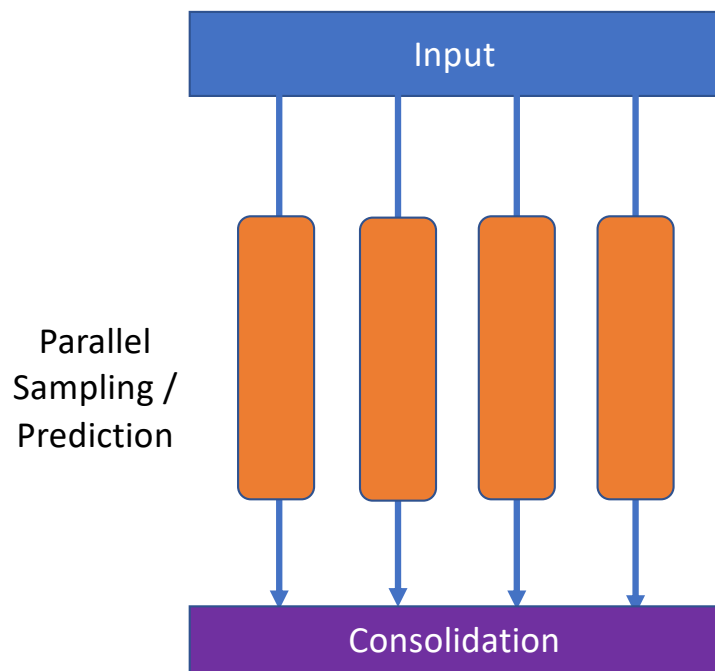
Imagine looking down from
the top of a cliff

What do you feel?

Minicolumns are plentiful in humans




What if minicolumns are predicting multiple futures?




- Perhaps we do parallel prediction in minicolumns
- Perhaps some of them can trigger the fear circuit (ref. the cliff scenario)
- Perhaps human thinking is not meant to be optimal, but use whatever works first
 - The more we sample, the more likely a solution close to optimal will occur
- **Question to ponder: How much performance can we get from increasing parallel search?**

Perhaps intelligence can be gained with multiple sampling?



Toward the quantification of cognition. Richard Granger. 2020.

Aiming a ball

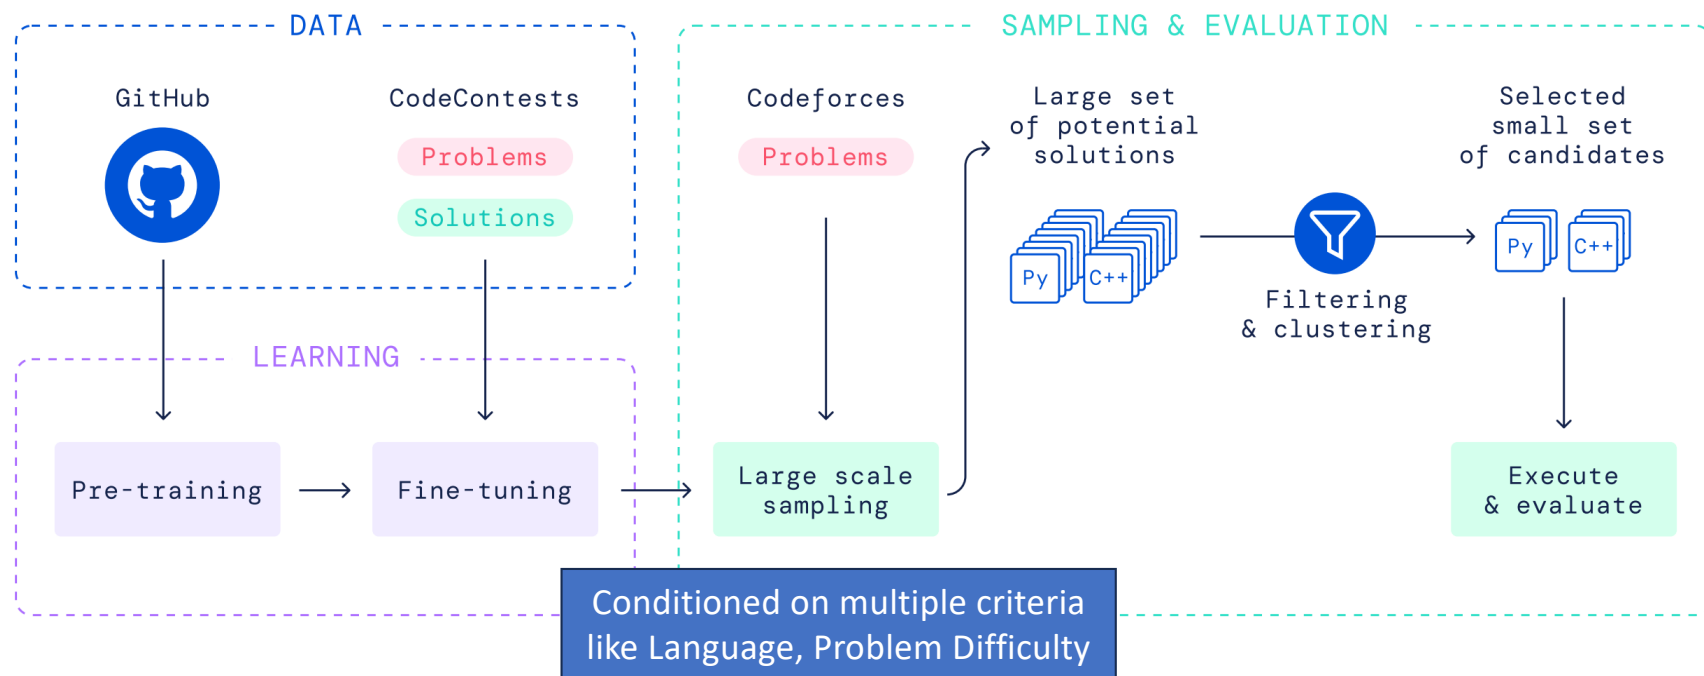


- Sampling is like throwing a ball
- Repeated throws have a higher chance to hit target (vary your starting angle etc.)
- **But what if the target is behind you?**

AlphaCode

Massive sampling and filtering to solve competitive programming questions

AlphaCode



Competition-Level Code Generation with AlphaCode. DeepMind. 2022

Encoder Input X:

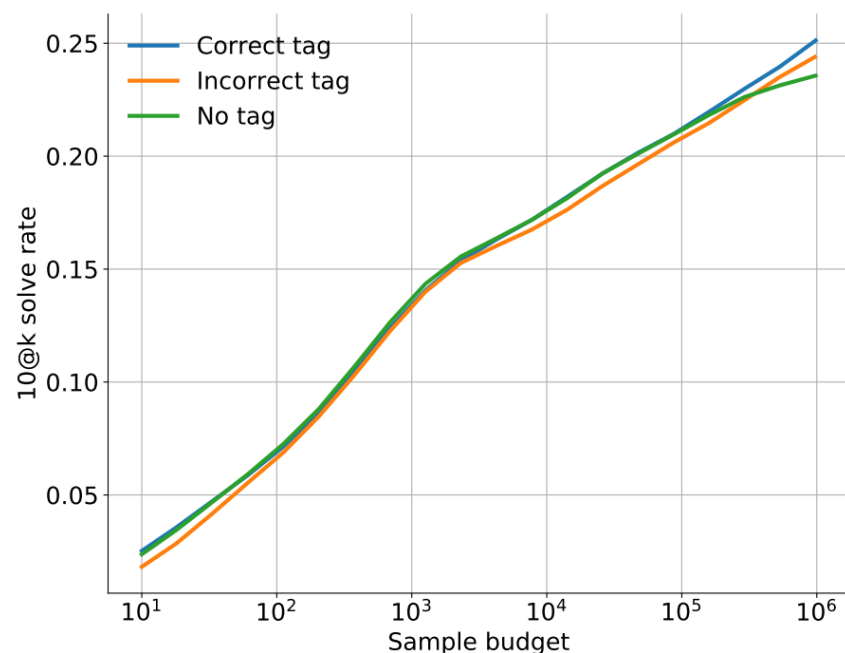
```
// RATING: 1200
// TAGS: math
// LANGUAGE IS cpp
// CORRECT SOLUTION
// n towns are arranged in a circle sequentially. The towns are numbered from 1
// to n in clockwise order. In the i-th town, there lives a singer with a
// repertoire of a_i minutes for each i ∈ [1, n].
```

Decoder Output Y:

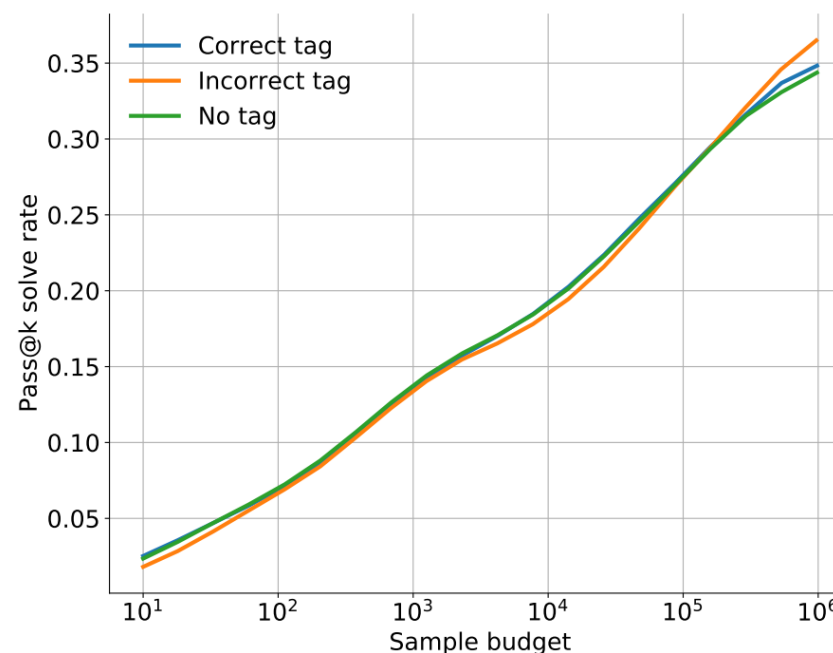
```
#include <bits/stdc++.h>
using namespace std;
void solve() {
    long long int n, i, j;
    cin >> n;
    vector<long long int> a(n);
    for (i = 0; i < n; i++) cin >> a[i];
    long long int sum = 0;
    for (i = 0; i < n; i++) sum += a[i];
    if (sum %
        cout << "NO" << '\n';
    return;
}
```

Competition-Level Code Generation with AlphaCode. DeepMind. 2022

Increased sampling = Increased Solve Rate



(a) 10 attempts per problem



(b) Unlimited attempts per problem


Appendix Figure A13 | **Conditioning on the CORRECT SOLUTION, the INCORRECT SOLUTION, or no tag.**

Competition-Level Code Generation with AlphaCode. DeepMind. 2022

Learning, Fast and Slow

Massive sampling and filtering to solve a dynamically changing maze

Two Networks – Fast and Slow



Neural Networks: Fast retrieval, slow learning

Memory: Slow retrieval, fast learning

Learning, Fast and Slow. John and Motani. 2023

Dynamic Environment




Figure 6: A sample maze environment of size 10x10. By default, the agent's start state is at the top left and the door is at the bottom right, but it can be varied. **(Left)** Obstacles before episode 50 form a vertical wall with a gap in the center across the mid-point. **(Right)** Obstacles after episode 50 form a horizontal wall with a gap in the center across the mid-point.

Increased sampling helps increase efficiency of solution

- Increased depth of sampling, increased breadth (threads) helps find shorter paths

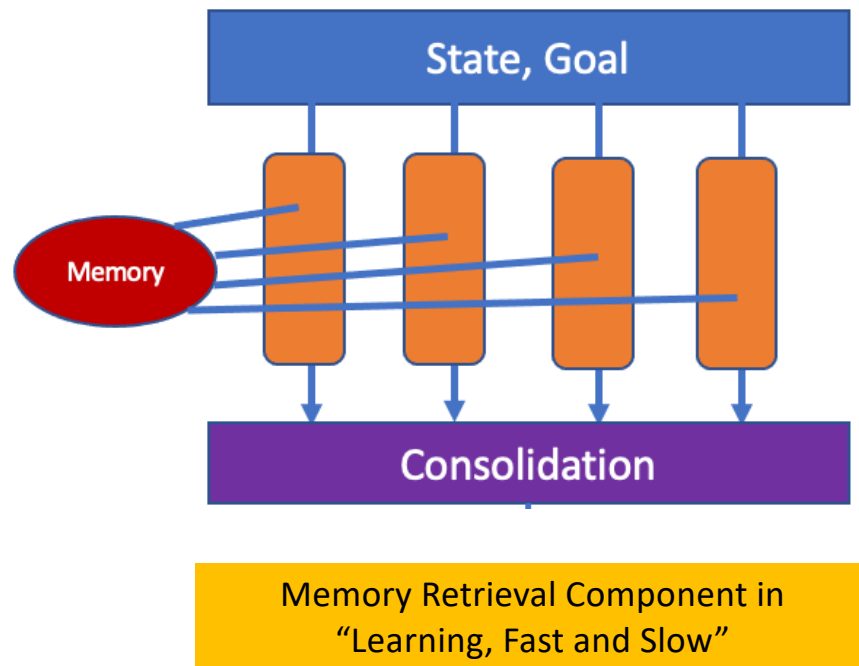


TABLE IX

ABLATION STUDY ON EFFICIENCY OF F&S AGENT ON A DYNAMIC 10x10 NAVIGATION TASK. LOWER IS BETTER (IN BOLD).

| Agent | Steps Above Minimum | | |
|--------------|----------------------|----------------------|-----------------------|
| | First 50 eps | Last 50 eps | Total |
| Baseline | 1029.5± 145.4 | 675.4± 223.3 | 1704.9± 280.6 |
| No Slow | 2625± 234.4 | 2517.0± 316.0 | 5142.7± 389.7 |
| No Fast | 2694.7± 216.8 | 2386.6± 445.0 | 5081.3± 496.2 |
| No Fast,Slow | 3890.6± 222.6 | 3853.0± 207.5 | 7743.6± 317.3 |
| 10 depth | 1225.5± 225.1 | 821.3± 292.3 | 2046.8± 455.8 |
| 50 depth | 941.2± 168.0 | 617.6± 115.4 | 1558.8± 250.6 |
| 50 threads | 1112.6± 216.0 | 761.2± 156.9 | 1873.8± 231.0 |
| 200 threads | 870.2 ± 152.9 | 521.2 ± 132.7 | 1391.4 ± 224.9 |


Learning, Fast and Slow. John and Motani. 2023

ARC Challenge


Massive sampling and filtering to solve image-based IQ puzzles

ARC Challenge


Example 1: **Input**




Example 1: **Output**




Example 2: **Input**




Example 2: **Output**




Example 3: **Input**



Example 3: **Output**



Test: **Input**



Test: **Output**





<https://lab42.global/arc/>

Overall Approach by Ryan Greenblatt

- Show the problem in image form, various text forms for the grid representation
 - The text representations include showing which **cells are occupied by different connected components of colors** and showing **diffs between the input and output** (in cases where the grid shapes are the same)
- Sample 8000 programs per problem (wow)
- Take most promising 12 programs (based on score on training examples) and revise outputs according to generated output vs actual training set output

The benefits of sampling



<https://redwoodresearch.substack.com/p/getting-50-sota-on-arc-agi-with-gpt>

The drawbacks of sampling

[–] **Aaron_Scher** 23d 

< || > ✕ 0 ✓

Nice! Do you have a sense of the total development (and run-time) cost of your solution? "Actually getting to 50% with this main idea took me about 6 days of work." I'm interested in the person-hours and API calls cost of this.

Reply




[–] **ryan_greenblatt** 22d 

< 14 > ✕ 3 ✓

I was pretty inefficient on iteration, so around \$40,000. Around 6 person days of development work, though a considerable amount of hours on each day.


Reply

👍 16 

<https://www.lesswrong.com/posts/Rdwui3wHxCeKb7feK/getting-50-sota-on-arc-agi-with-gpt-4o#XheMa2FNjjdQGSqJ>

My solution – sample with agents with different views

- Each agent views the grid differently, like grid, object, or pixel level
- Generate multiple potential Python programs
- Filter solutions by those which pass training examples
- 20 views * 3 samples = 60 for each problem
- **50 out of 111 public training set tasks solved (45%)**



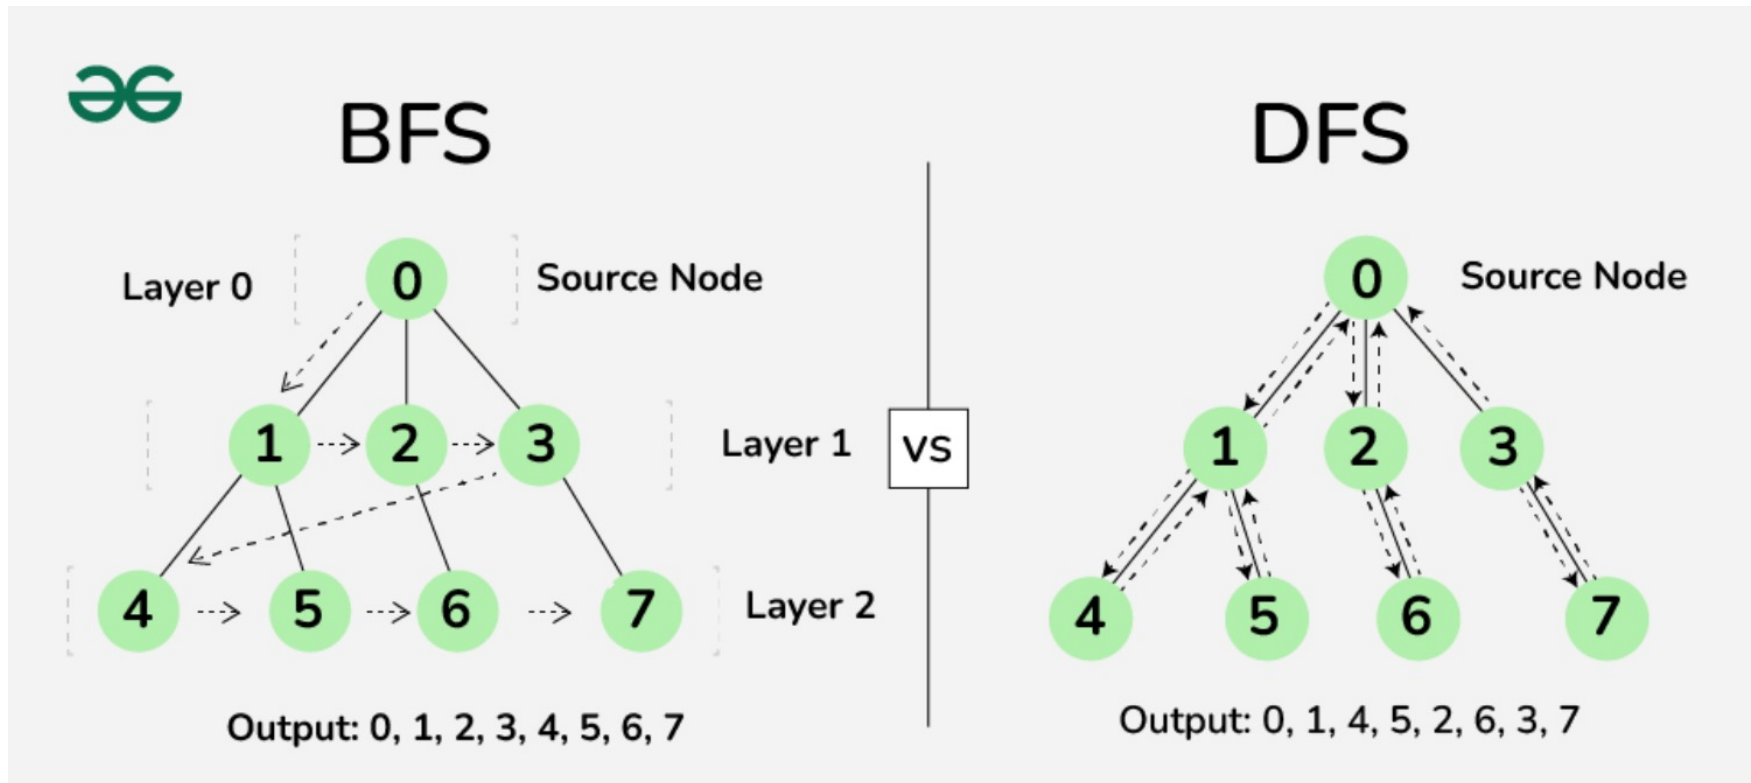
LLMs as a System of Multiple Expert Agents. John and Motani. 2023.

More Efficient Sampling

Sample breadth-wise, sample depth-wise

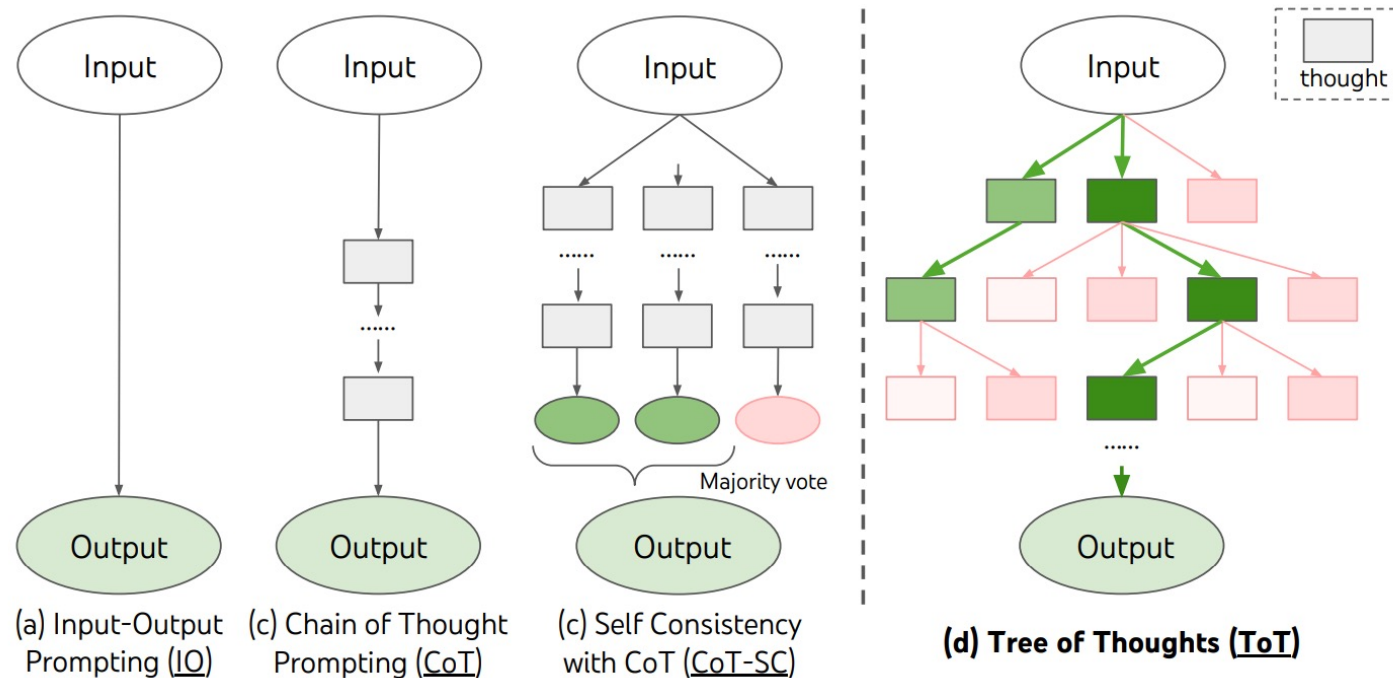
Sample Monte-Carlo Tree Search-wise

Breadth-First Search vs Depth-First Search



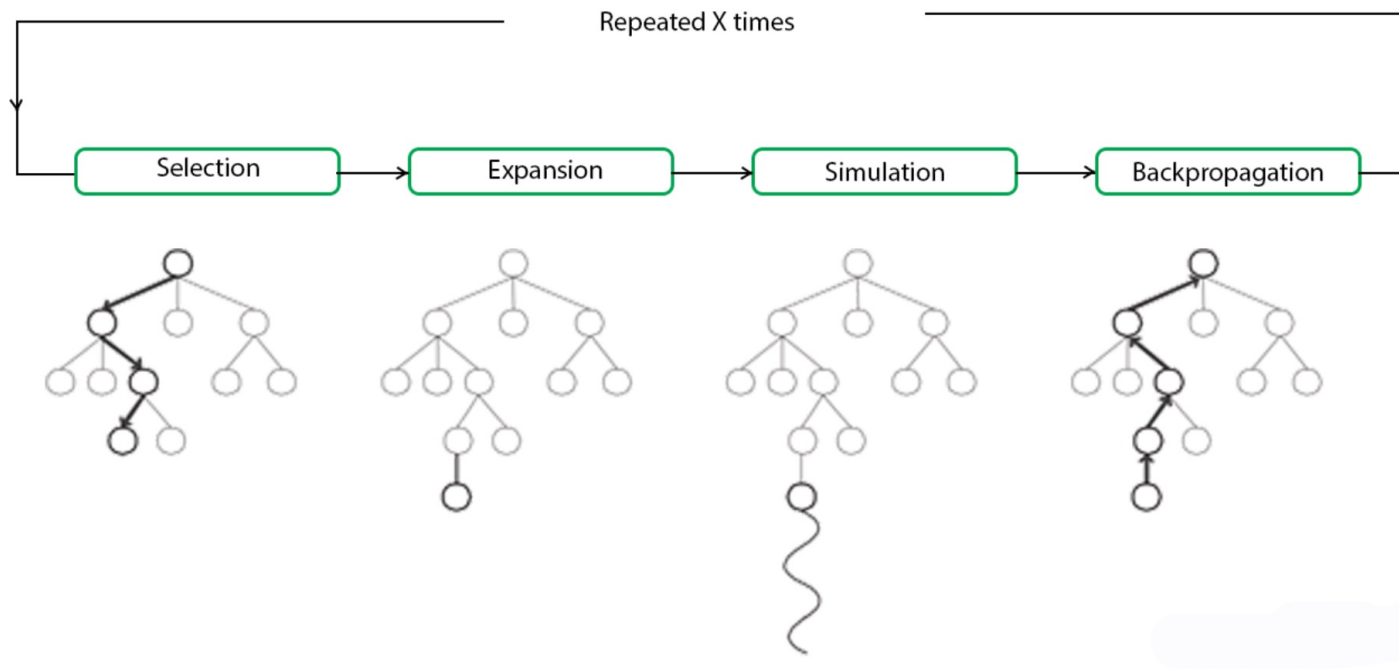
<https://www.geeksforgeeks.org/difference-between-bfs-and-dfs/>

Tree of Thought: Expand out most promising node, prune breadth-wise or depth-wise




Tree of Thoughts: Deliberate Problem Solving with Large Language Models. Yao et al. 2023.

Monte-Carlo Tree Search: Search with Explore and Exploit



<https://www.geeksforgeeks.org/ml-monte-carlo-tree-search-mcts/>

MCTS can be an efficient way to search and improve



AlphaZero: MCTS used to improve policy and value networks

| Datasets | Zero-Shot CoT | One-turn Self-refine | 4-rollouts MCTSr | 8-rollouts MCTSr | Example Nums |
|----------|---------------|----------------------|------------------|------------------|--------------|
| GSM8K | 977 | 1147 | 1227 | 1275 | 1319 |
| | 74.07% | 86.96% | 93.03% | 96.66% | |
| GSM-Hard | 336 | 440 | 526 | 600 | 1319 |
| | 25.47% | 33.36% | 39.88% | 45.49% | |

Table 1: Performance of MCTSr on the GSM Dataset

MCTSr: Using Llama3 8B to evaluate promising states to expand with a refinement step

Accessing GPT-4 level Mathematical Olympiad Solutions via Monte Carlo Tree Self-refine with LLaMa-3 8B: A Technical Report. Zhang et al. 2024.

My thoughts

- More complicated forms of sampling like BFS, DFS, MCTS require keeping track of existing states globally
- Minicolumns in the brain may not have the infrastructure to do that and may just be doing parallel search with different biases
- BFS, DFS, MCTS are traditionally sequential operations and VERY SLOW to implement
 - Potentially able to do asynchronous operations at the cost of some non-optimality (ref. AlphaGo setting node value to huge negative number to prevent exploitation when one branch is being calculated)

Questions to Ponder

- Can sampling help if model does not know how to generate the answer?
- How do we bias the generation effectively?
- Is System 2 thinking simply just more deliberate use of the parallel generation mechanism?
- Are more complex forms of search, like BFS or DFS or MCTS, needed?